


# An Approach to Routing in a Clos


Randy Bush <randy@psg.com>

Arrcus & IIJ Research

# This Works




# This Might Work



# This Won't Work



# This Works (Clos Network)



# Clos is Not an Acronym

*Clos, Charles (Mar 1953)*

*"A study of non-blocking switching networks"*

Bell System Technical Journal. 32  
(2): 406-424

For Example:  
IIJ Built a Second  
Medium Scale Data  
Center (MSDC)  
in Shiroyi/Chiba  
Capacity of 6k Racks

How Do You Route  
In Something of  
This Scale?



OSPF OK to 500 Nodes

IS-IS good to 1,000

Limited Because They  
Repeatedly Flood  
Everything


# Your Clos on IS-IS or OSPF



BGP Scales Because  
It Signals  
Only Changes


So BGP has become  
common in MSDCs

# BGP is Quiet as Updates are Infrequent



# ECMP can be Very Wide


## 32, 64, even 128




BTW, Every Rack  
is (often) an AS

Get Over It

# But What is the Decision Process?



# Do You Want to Write BGP Policy for Massive ECMP?





# Consult the Professor



**Edsger W Dijkstra**  
**1930-2002**

## Shortest Path First

# BGP-SPF



The Path Calculation of IS-IS  
With the Update Rate of BGP

SPF?

I thought BGP was path  
vector, not link state!

# s/Best Path/SPF/


- New SAFI
- NLRI format exactly same as BGP LS (RFC 7752) Address Family to carry link state information
- BGP runs Dijkstra instead of Best Path Decision process
- BGP MP (new SAFI) and BGP-LS Node attribute for compatibility
- Peering Models: eBGP, iBGP, RR

# BGP4 Classic

**Neighbor  
Distribution  
Route Reflection  
Outbound Policy**

**AS-Path Length  
EGP vs IGP  
Arrival Order  
Non-deterministic  
MED  
IGP metric  
Tie Break**

**Inbound Policy  
Link State**



BGP-  
SPF

AS-Path Length  
EGP vs IGP  
Arrival Order  
**Removed!**  
Non deterministic  
MED  
IGP metric  
Tie Break

# BGP-SPF

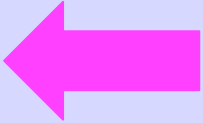
- Next-Hop and Path Attributes come for free with BGP Link-State Address Family
  - Needed for RFC 4271 error handling
- Decision Process Phases 1 and 2 (best path) replaced by SPF algorithm (AKA Dijkstra)
- Decision Process Phase 3 (tie break) may be skipped as NLRI is unique per BGP speaker
- Need to assure the most recent version of NLRI is always used and re-advertised
  - Augmented with sequence numbers

# BGP-SPF

- Starting with greatly simplified SPF with P2P only links in single area (i.e., SPT)
- Should scale very well to many use cases
- Could support computation of LFAs, Segment Routing SIDs, and other IGP features
  - BGP-LS format includes necessary Link-State
- Link-State AF is dual-stack AF since both IPv4 and IPv6 addresses/prefixes advertised
  - BGP-LS format also supports VPNs but SPF behavior not defined
  - Work needed to define interaction with existing unicast AFs
    - Matter of local implementation policy



# Peering Model

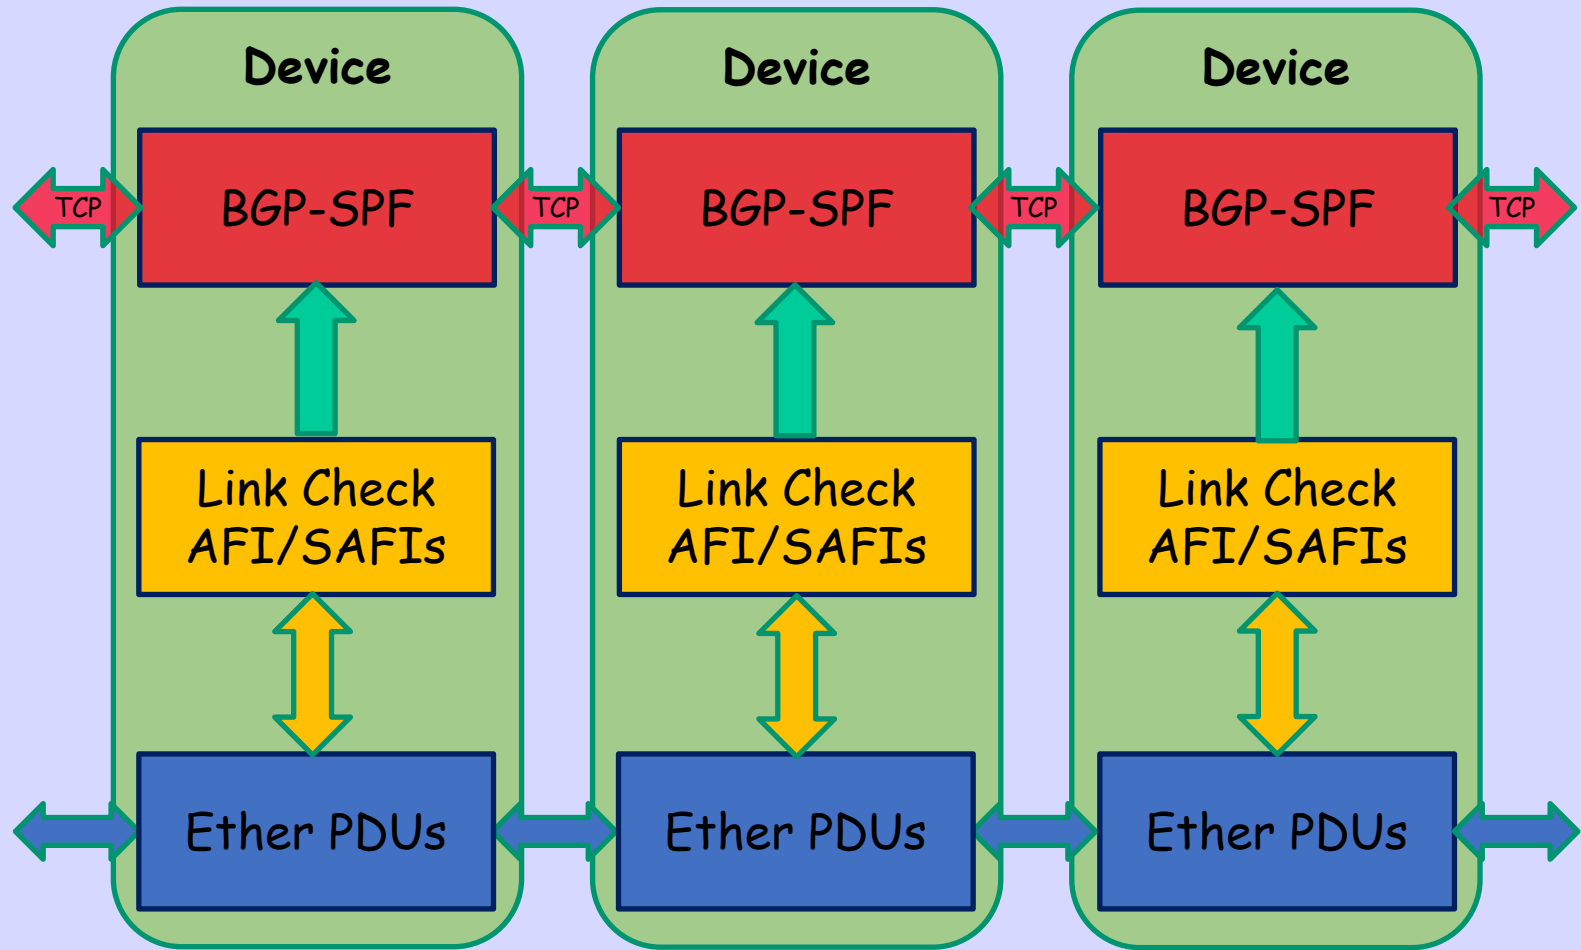
- BGP sessions, optionally with Route-Reflector or controller hierarchy
  - Link discovery/liveliness detection outside of BGP 
- RR hierarchy can be less than fully connected but must provide redundancy
  - Must not be dependent on SPF for connectivity
- Controller could learn the expected topology through some other means and inject it
  - SPF Computation is distributed though
  - Similar to "Jupiter Rising: A Decade of Clos Topologies and Centralized Control in Google's Datacenter Network"

How Does BGP-SPF  
Learn Link State  
so it can Build  
the Topology?

# Motivation

- BGP-SPF needs link neighbor discovery, liveness, and addressability
- LLDP is an IEEE protocol, complex, and 'hard' (IPR) to extend past 1500 bytes
- We wanted something simple and saw no real need for the complexities of CLNP, ...
- So we propose a new EtherType with TLVs
- We discuss Ether payloads, not framing

# Topology / Routing Stack




MAC Link State exchanged over raw Ethernet and pushed up stack  
Add the AFI/SAFI data IP-Level Liveness Check  
BGP-SPF uses link data to discover and build the topology database

```

0      1      2      3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| Version  | Transmission Sequence Number |L| Datagram ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| Number   | Datagram Length              |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|                               Checksum   |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|                               Payload...  |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

```



```

0      1      2      3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| PDU Type | Payload Length              ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~          | Payload ...                 |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| Sig Type | Signature Length            | ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~          Signature                      ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

```

# Why not Just Use TCP?

- When this runs, there are no IP Addresses
- This protocol is to Learn IP Addresses
- So it is a cheap TCP-like protocol
- Reassembly of out of order Datagrams
- Retransmission with Back-off
- PDUs are ACKnowledged
- ...

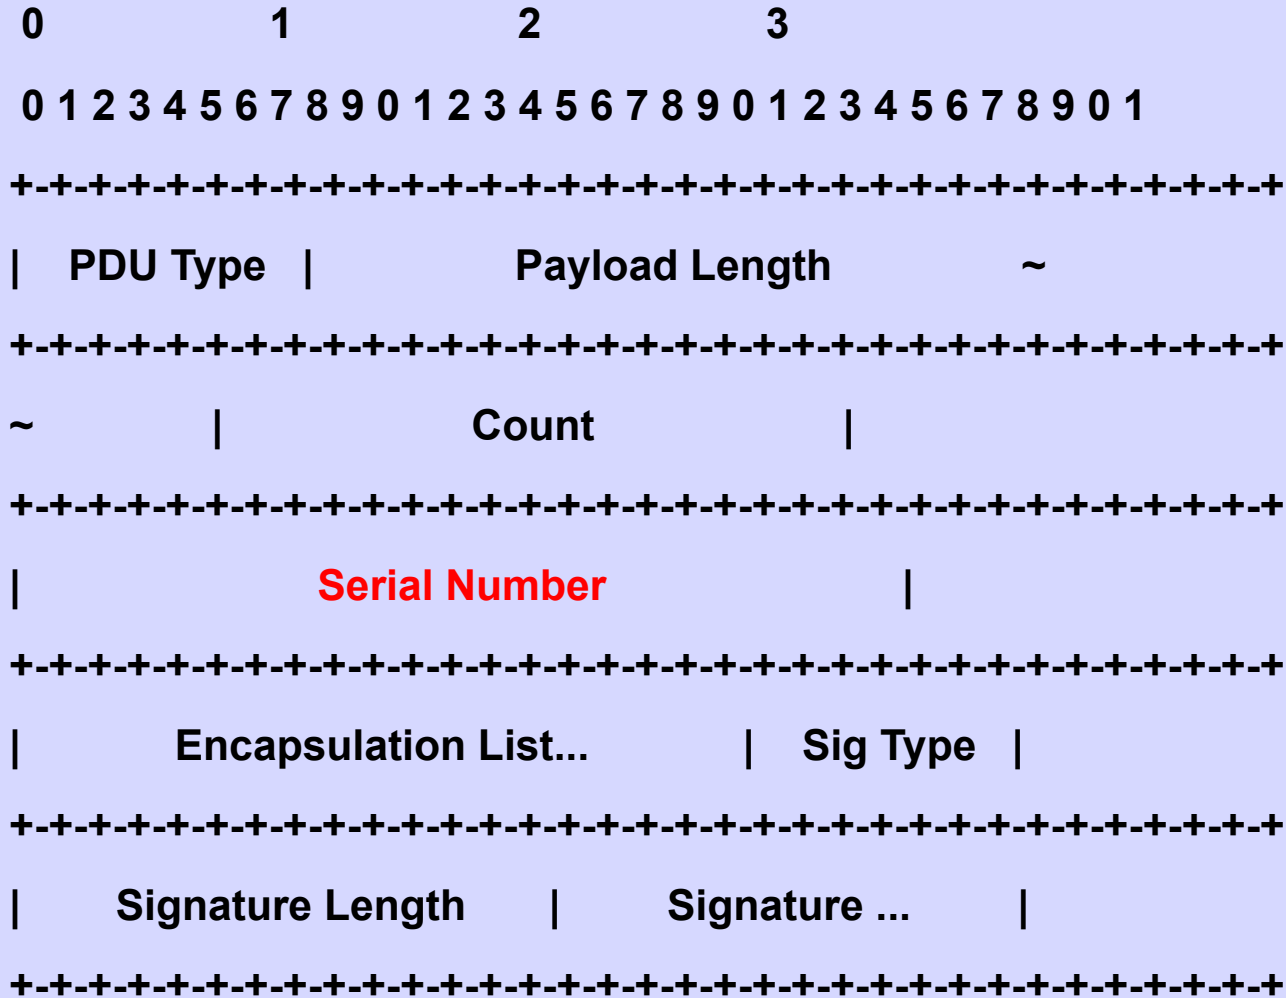
# Fully Stateful Session Per Peer

# Graceful Restart

## State May Be Resumed à la BGP




# Encaps etc PDUs



# OPEN PDU

```
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| PDU Type = 1 |          Payload Length          ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~          |          Nonce          ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~          | LLEI Length |      My LLEI      |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+~
~          | AttrCount |          ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~  Attribute List ... | Auth Type | Key Length ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~          |          Key ...          |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|          Serial Number          |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| Sig Type |      Signature Length      | Signature ... |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

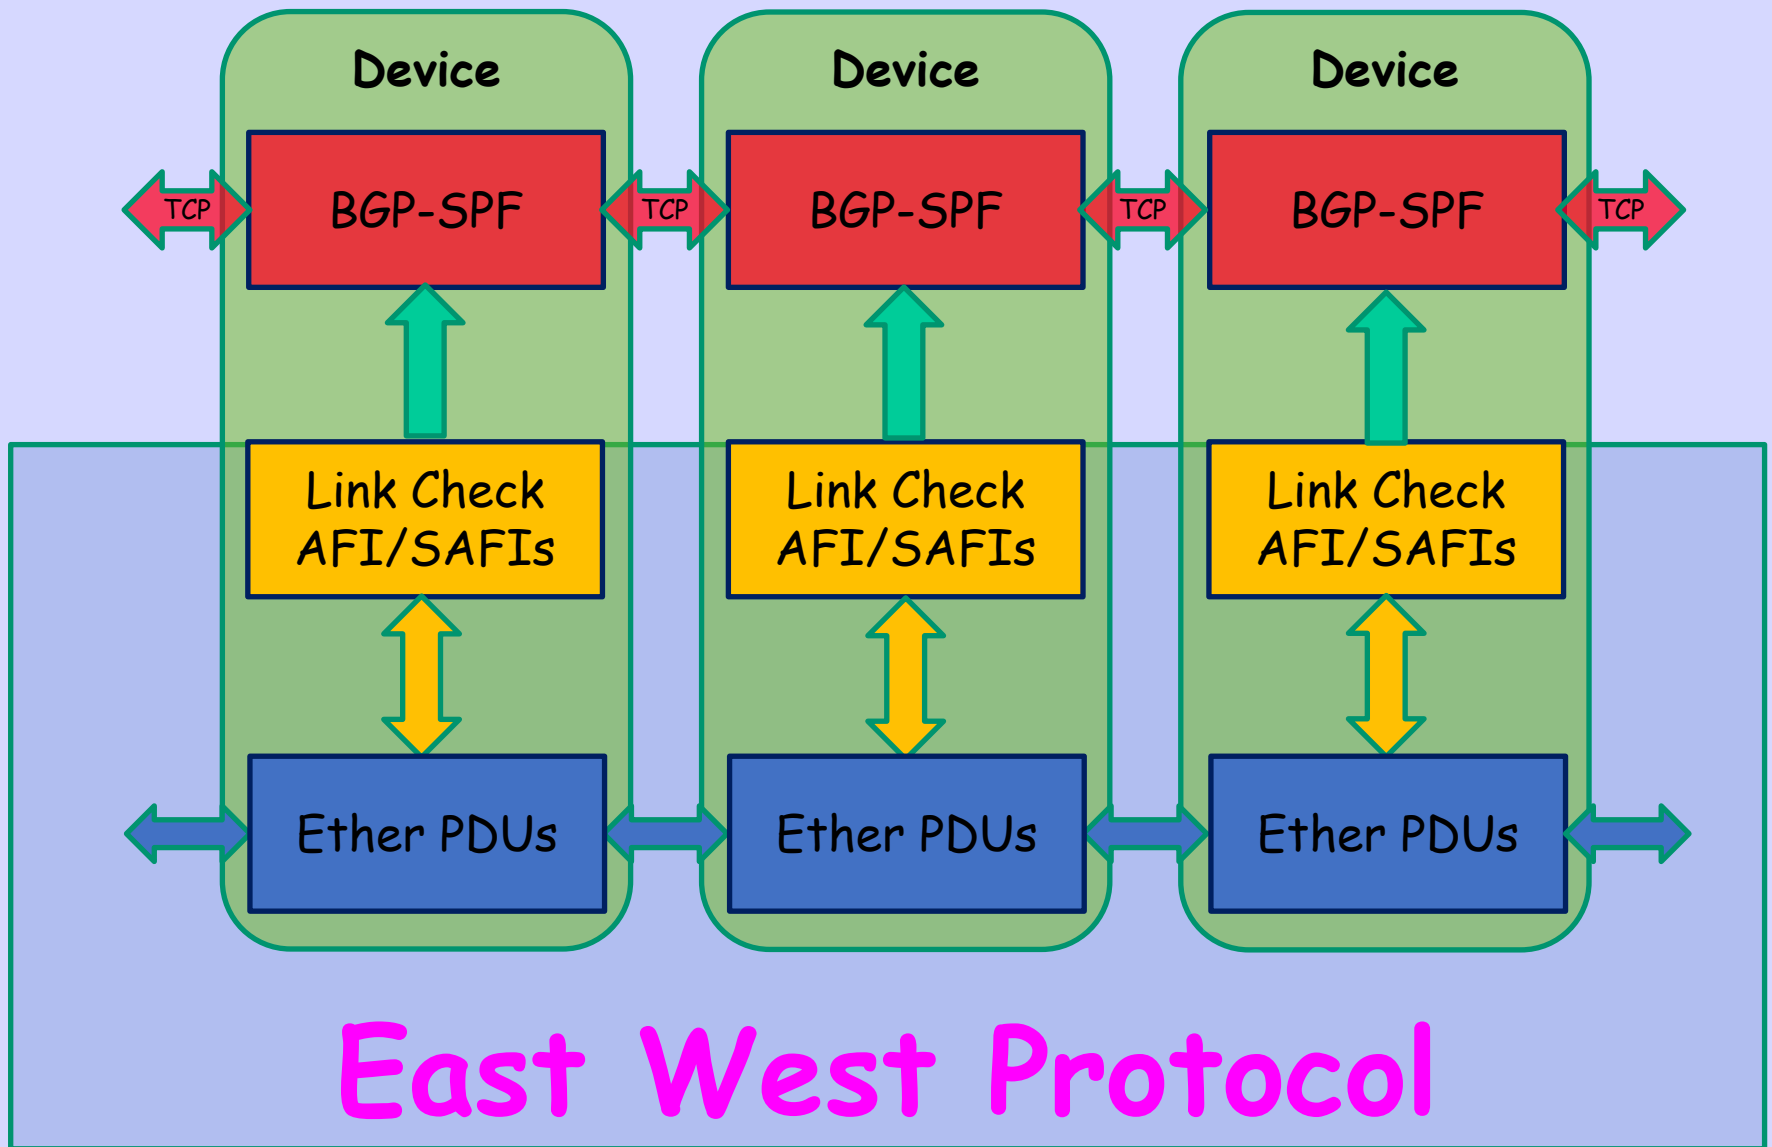
# Announce/Withdraw



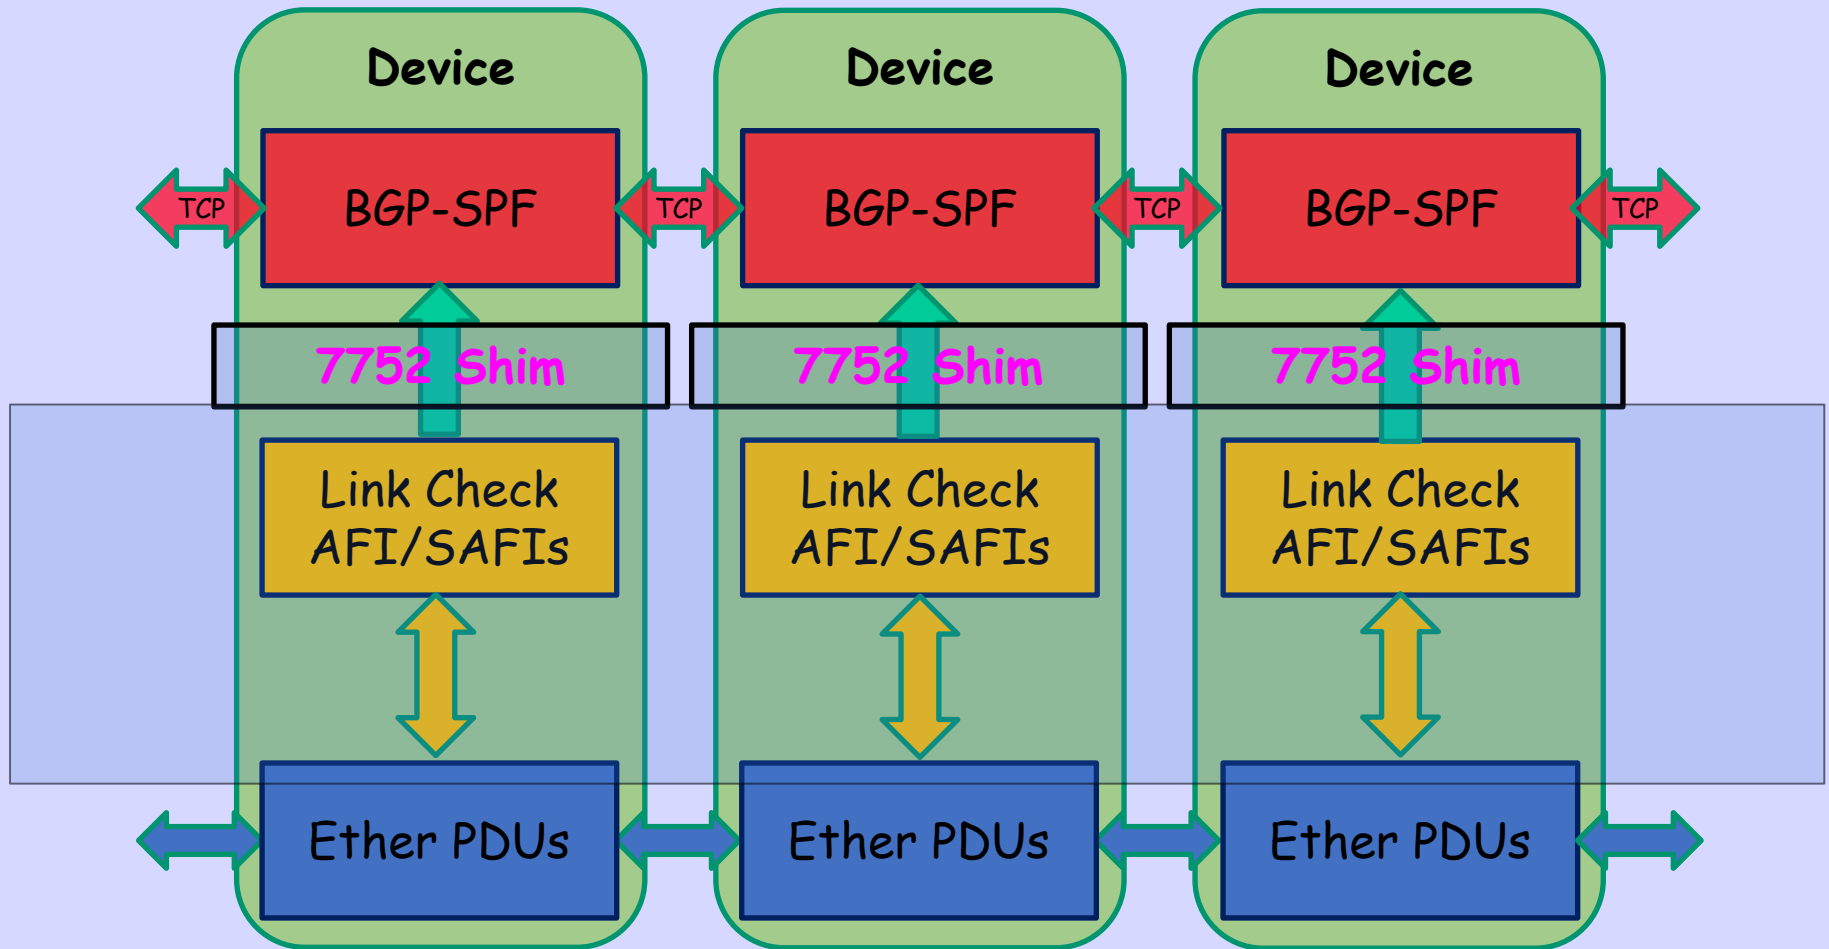
# Explicit ACK/ERROR

[illegible]

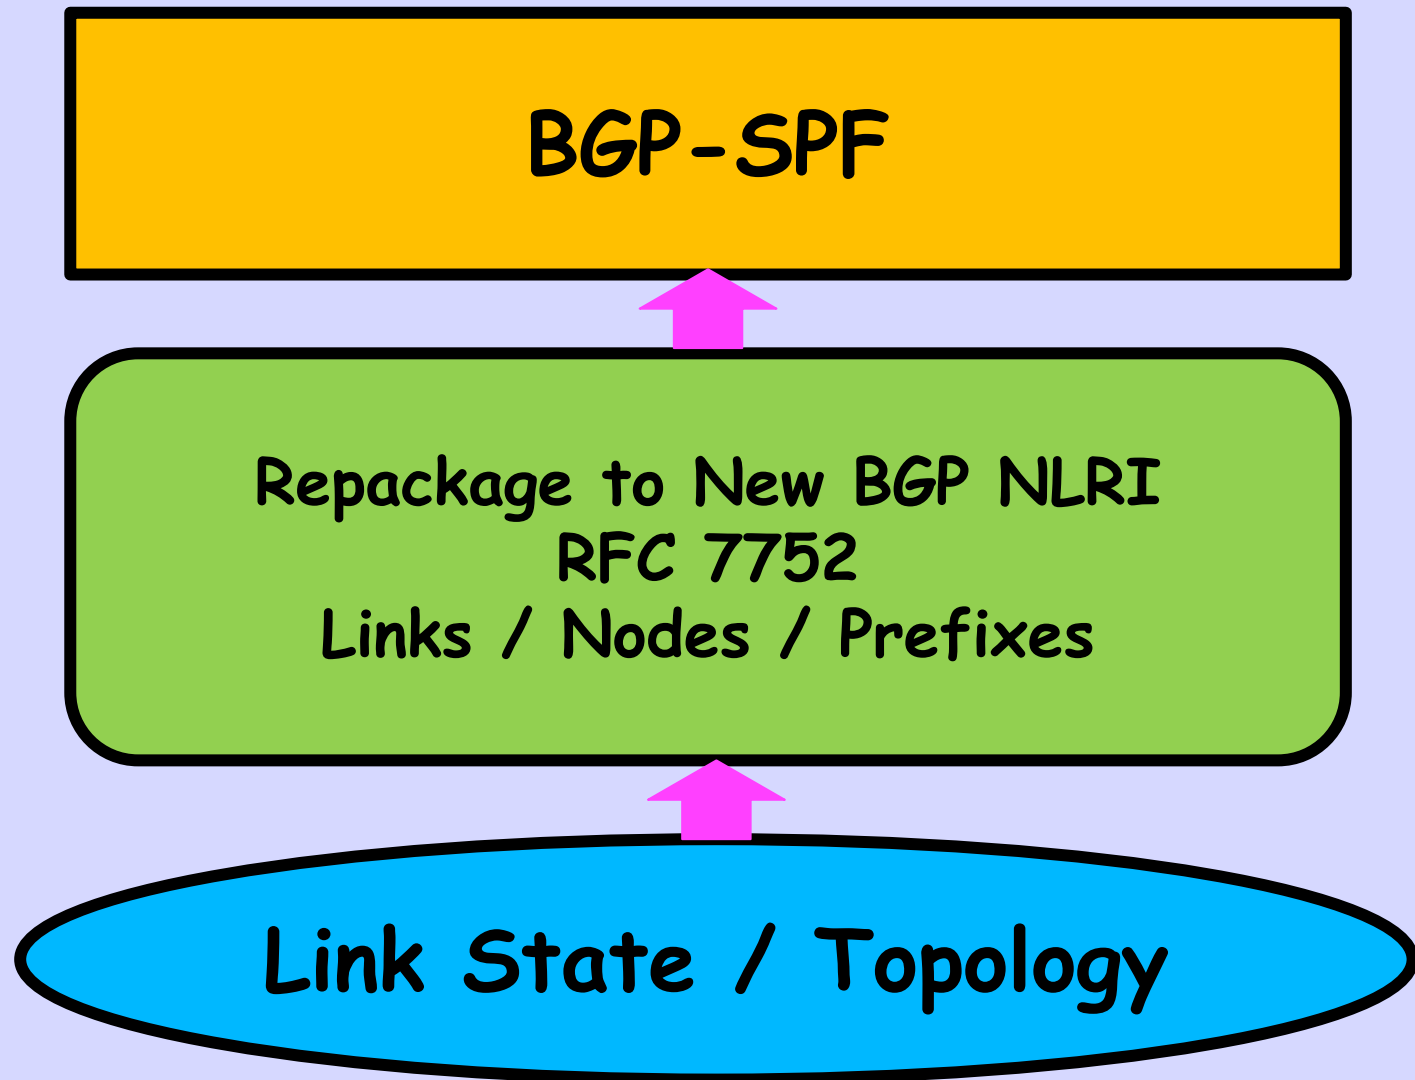
```
EType, Error Code, Error Hint = 0    /* no error, just an ACK */
```



# North/South Protocol



# BGP-LS for BGP-SPF



# How Does BGP-SPF Start?

- For BGP-SPF to build topology and state, need to peer with BGP-SPF neighbors
- But we do not want to configure it more than necessary
- Why not extend L3DL to pass BGP config data to my peer?

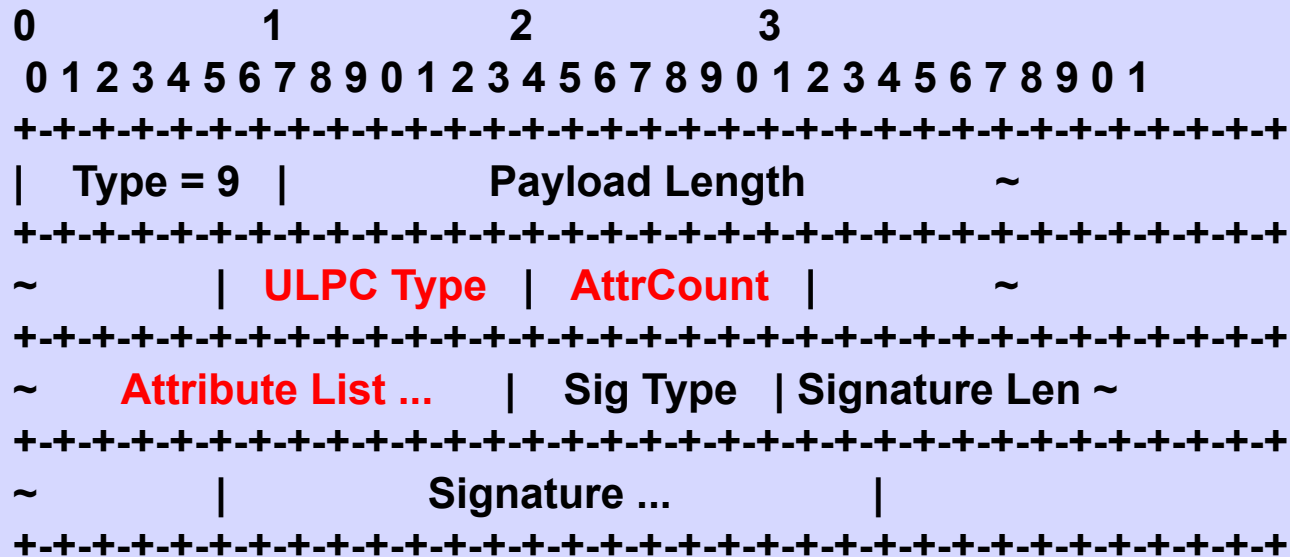


L3DL-ULPC

Upper Layer Protocol  
Configuration

**draft-ymbk-lsvr-l3dl-ulpc**

# L3DL PDU for ULPC



Provide the minimal set  
of configuration  
parameters for BGP  
OPEN to succeed

Not to replace or  
conflict with data  
exchanged by  
BGP OPEN

Multiple sources of  
truth are a recipe for  
complexity and pain

# ULPC for BGP

```
0          1          2          3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| Attr Type = 1 | Attr Len = 48 |      My ASN      ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~
|
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

```
0          1          2          3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| Attr Type = 2 | Attr Len = 56 |  My IPv4 Peering Address  ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~
| Prefix Len |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

```
0          1          2          3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
| Attr Type = 4 | Attr Len |                                     ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
~
| BGP Authentication Data ...                                     ~
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

BGP Can Now OPEN

And that will exchange  
all the rest

# There is Running Code

- Open Source Python3 for LSOE, an early version of L3DL
- GoLang Source for current L3DL. We hope to be allowed to open source



BTW, There is  
No IPR in these  
Standards Proposals